

THE NO-CODE AI BUILDER'S PLAYBOOK

From Idea to *Shipped*.

The step-by-step path for PMs, founders, and operators who want to ship a real AI product with no code. Validate the idea, build with Lovable and Supabase, wire the AI, and launch, with the one gate that tells you each stage is actually done.

FORMAT

Playbook

STAGES

5

READ TIME

12 min

By Aki Wijesundara, PhD & Manu Jayawardana

Founders, Snapdrum · Instructors at Maven

Anyone can start a build.

Finishing is the rare part.

No-code tools made starting an AI product almost free. You can have a screen by lunch. The hard part was never the first screen. It is getting from a screen that demos to a product someone actually uses, trusts, and comes back to.

Most no-code AI products die in that gap, and the tool is rarely the reason. They die because the builder built before they validated, scoped to everything instead of one job, or shipped without ever testing it works for someone other than themselves.

The five stages below map to the way a product actually gets shipped, roughly in the order the mistakes tend to arrive. Each stage carries a checklist you can run against your own idea today, plus the single gate that means the stage is finished.

How to use this: keep it open while you build. Run each stage's checklist before you move on. When the build stalls, come back and find the stage you rushed. It is almost always one of these.

NO.	STAGE	THE GATE THAT MEANS DONE
01	Validate the idea	A real person with the problem said they would use it, in their words.
02	Shape the build	The whole product is scoped to one job and written down as a brief.
03	Build the core	A working app with real data and accounts, live on a real URL.
04	Add the AI and automation	The AI does one job well and can never run away with cost.
05	Evaluate and ship	It works for someone who is not you, and you can roll it back.

Earn the right to *build*.

The cheapest product to fix is the one you have not built yet. Before a single screen, prove the problem is real and worth solving.

Find the real *problem*

- ◆ **PAINFUL** You can name one specific person and the exact moment the problem bites them. "Everyone" is not a customer.
 - ◆ **TODAY** They already solve it with a clumsy workaround: a spreadsheet, copy-paste, a manual back-and-forth. The pain exists, you are not inventing it.
 - ◆ **WORTH IT** Solving it saves real time or money, enough that they would change their routine to use what you make.
-

Talk before you *build*

- ◆ **FIVE VOICES** You have spoken to at least five people with the problem and heard the same pain in their words, not yours.
 - ◆ **DEMAND** At least one said a version of "I would use that today." Polite interest is not demand, and a compliment is not a commitment.
-

Pick the *wedge*

- ◆ **ONE JOB** You can describe what the product does in one sentence: one job, one user. The roadmap can wait.
- ◆ **WHY AI** AI earns its place here, because it reads, writes, sorts, or decides at a scale a person cannot. It is not bolted on for the pitch.

Turn the idea into a *spec*.

A vague idea becomes an endless build. Before you open a tool, write down exactly what you are making, for whom, and what done looks like.

Scope to *one* thing

- ◆ **ONE WORKFLOW** The first version does one workflow end to end, not three workflows halfway. Depth beats breadth on a first ship.
 - ◆ **DONE DEFINED** You have written what "working" means in plain language, the way a user would judge it, not "it feels good."
-

Write the *brief*

- ◆ **ONE PAGE** Who it is for, the one job, the screens, the data it stores, and what the AI does. This is the prompt you will hand the no-code tool.
- ◆ **IN AND OUT** For the AI step, you have written what goes in and what good output looks like, with two real examples to aim at.

SHIP IN V1

The one job, done properly.

- The single workflow, end to end
- Real accounts and saved data
- One AI step that clearly helps
- The three obvious error cases handled

PARK FOR V2

Everything that can wait.

- The second and third workflows
- Settings, themes, admin panels
- Integrations nobody has asked for
- Polish on screens no user reaches yet

Build the product with *no code*.

Now you open the tools. The goal of this stage is a live URL: a real app, with real accounts and real data, that you could put in front of a user today.

The app and the *data*

- ◆ **LIVE URL** Lovable builds the screens and the flow from your brief. The product has a real address on the internet, not a demo that only runs on your laptop.
- ◆ **REAL BACKEND** Supabase holds the data, the accounts, and any uploads. Nothing important lives only in the browser or a single session.
- ◆ **IT SAVES** A user can close the tab, come back tomorrow, and their work is still there. State survives a refresh.

When no-code hits a *wall*

- ◆ **EXTEND** When the visual builder cannot express something, Claude Code or Cursor adds the missing logic. You extend the app, you do not rebuild it from scratch.
- ◆ **SMALL STEPS** Each change is small and you watch it work before the next one. You never paste a hundred lines you cannot read and hope.

TOOL	WHAT IT DOES	REACH FOR IT WHEN
Lovable	Builds the app and screens from a prompt	You need a working web product fast, with no front-end code.
Supabase	The database, accounts, and file storage	Your product needs logins, saved data, or uploads.
Claude Code, Cursor, Cowork	Your AI pair to build and extend the app	The no-code builder cannot express something you need.
n8n	Automates the steps between tools	Work should happen on a trigger or schedule, with no human.
Hermes	A persistent agent that does real work in chat	The product needs an agent that acts on its own, in WhatsApp or Slack.

Add the AI, and keep it on a *leash*.

The AI is one ingredient, not the whole dish. Give it one clear job, check what it returns, and make sure it can never quietly run up a bill or do harm.

Give the AI *one job*

- ◆ **ONE JOB** The AI does one thing in your product: it drafts, sorts, summarizes, or answers. One clear job beats a vague assistant that does everything badly.
- ◆ **REAL PROMPT** The instruction that drives it lives in your build as a saved, editable prompt, not a sentence you typed once and forgot. You can change it and see the effect.
- ◆ **CHECKED** You see what the AI returns before a user does. A wrong or empty answer shows a graceful fallback, it does not break the screen.

Wire it to *act*

- ◆ **TRIGGERS** n8n connects the steps: a new signup, a form, a schedule. The boring work happens without you in the loop.
- ◆ **STAND IT UP** When the job is ongoing rather than a one-off, a persistent agent like Hermes runs it and meets users where they are, in WhatsApp or Slack. You deploy it once and it keeps working.
- ◆ **SAFE DEFAULT** Anything that sends messages, charges money, or deletes data is tested on yourself first and can be turned off with one switch.

Watch the *cost*

- ◆ **A CEILING** Every AI call has a limit on length and a sensible model behind it. You know roughly what one user costs, and a cap stops a loop from draining the account.

ONE GOOD STEP FIRST

Reach for one reliable AI step with a clear job before you wire five together. Most products do not need a swarm of agents, they need one step that earns trust. **Add the second only when the first is solid and a user is asking for it.**

Prove it works for *someone else*.

Your product works for you, the person who built it and knows exactly what to type. Real use is everyone else. TRACE is the loop that turns real sessions into a product you can trust to ship.

T	Trace	Capture every real session: what the user typed, what the AI returned, what they did next. You cannot improve what you did not record.
R	Read	Read real sessions by hand, not just the happy path you demoed. Watch where users get confused or quietly give up.
A	Analyze	Group what goes wrong into a few common patterns. Fix the one that bites the most people, not the loudest single complaint.
C	Codify	Write down what a good result looks like as a simple test you can run again. Now "good" is something you check, not a feeling.
E	Enforce	Run that test before every change goes live. A change that breaks what used to work does not ship.

Before it goes *live*

- ◆ **ONE REAL USER** At least one person who is not you has used it start to finish and got value, without you sitting next to them.
- ◆ **REVERSIBLE** You can undo the last change in minutes, and there is an off switch if something goes wrong in front of a real user.
- ◆ **YOU WILL KNOW** Basic tracking is on, so you can see who used it, what worked, and what broke, without guessing.

Ship-readiness, *one page.*

NO.	STAGE	THE GATE THAT MEANS READY
01	Validate the idea	A real person with the problem said, in their words, they would use it.
02	Shape the build	Scoped to one workflow, written as a one-page brief, with done defined.
03	Build the core	Live URL, real accounts and data, survives a refresh. No-code first, extend only at the wall.
04	Add the AI and automation	One clear AI job, output checked, automations safe, a hard cost cap.
05	Evaluate and ship	TRACE closed, one real user got value, reversible with a kill switch.

RED FLAGS: YOU ARE NOT READY IF

You have not watched anyone but yourself use it. The product is three half-built features instead of one finished one. The AI can run up cost with no ceiling. An automation can send or delete without a test. You cannot undo your last change in under a minute.

The pattern underneath all five: earn the right to build, scope to one job, get it real, give the AI a leash, and prove it works for someone else before you ship. Demos skip these. Products are made of them.

Run this against *your* build.

Print this, or keep it next to the screen while you build. Tick a box only when it is true in the product, not in your head.

Idea and *scope*

- ☐ **VALIDATED** A real person with the problem told me, in their words, they would use it.
- ☐ **ONE JOB** I can say what the product does in one sentence: one user, one workflow.

Build and *AI*

- ☐ **REAL** There is a live URL with real accounts and data that survives a refresh.
- ☐ **LEASH** The AI has one job, its output is checked, and there is a hard cost cap.

Ship and *trust*

- ☐ **TESTED** Someone who is not me used it end to end and got value on their own.
- ☐ **REVERSIBLE** I can undo my last change in minutes, and there is a kill switch.

BUILD AND SHIP THE WHOLE THING WITH US

Agentic AI Builder's Bootcamp · For PMs, Founders, Operators

Live with Aki and Manu. Validate, build with no code, and ship a real AI product, then present it on Demo Day. Office hours, recordings, lifetime access.

ENROLL